

# Project's Write Up

## System Explanation

The AI Path Planner demonstrates grid-based navigation using classical AI search algorithms.


**Input:** 20×20 grid with random obstacles.

**Output:** Shortest path, path cost, execution time, and nodes explored.

**Techniques:** BFS, UCS, A\*(Manhattan), A\*(Euclidean), Goal-Conditioned Q-Learning (Reinforcement Learning)

### **Pipeline:**

Grid Generation → Algorithm Execution → Visualization → Performance Metrics



**Limitations:** RL training currently limited to small grids and discrete actions  
A\* operates on static maps only (no dynamic obstacle re-planning)

**Mitigation:** Code modularized for future enhancements (dynamic path re-planning, weighted terrains).

## Feature Table

Description	Platform	Complete (1-5)	Code	Notes
<b>Grid generation</b>	Local	5	Python	Fully functional with seed option
<b>BFS</b>	Local	5	Python	Guaranteed optimal path
<b>UCS</b>	Local	5	Python	Handles uniform edge cost
<b>A* (Manhattan)</b>	Local	5	Python	Most efficient and fastest

<b>A* (Euclidean)</b>	Local	4.5	Python	Slightly slower due to float operations
<b>Visualization (4-panel)</b>	Local	5	Matplotlib	Animated exploration comparison
<b>Performance metrics &amp; logging</b>	Local	5	JSON / Matplotlib	Generates charts automatically
<b>Goal-Conditioned Q-Learning</b>	Local	5	Python	Goal-Conditioned Q-Learning
<b>RL vs A* Comparison</b>	Local	5	Python / Matplotlib	Displays both policies side-by-side

## External Tools & Libraries

- **NumPy** : Grid representation and matrix operations
- **Matplotlib** : Visualization and charting
- **JSON** : Result logging
- **OS** :File system interaction

No external datasets used; all code is original and self-contained.


## Results & Analysis

- All algorithms found the optimal path.
  - **A\*(Manhattan)** achieved the best efficiency with 54.9% fewer explored nodes than BFS and fastest runtime (3.38 s).
  - BFS and UCS showed similar node counts but higher execution times.
  - A\*(Euclidean) was accurate but slightly slower due to floating-point distance computations.
- These results confirm that heuristic-guided search significantly reduces search space while preserving optimality.

## ***Reinforcement Learning***

- Introduced a **Goal-Conditioned Q-Learning agent** trained on user-selected goals via interactive grid clicks.
- Agent trained over **2400 episodes** (4 segments × 600 episodes each) and achieved **~100 % segment success** by convergence.

- During evaluation, both **RL** and **A\*** completed the multi-goal route:



### Interpretation:

The RL agent learns to generalize goal-directed behavior through experience, trading off minimal path optimality for **significant gains in efficiency**.

This experiment highlights how **heuristic search (A\*)** and **reinforcement learning** can complement each other: A\* ensures optimality, while RL enables faster, adaptive navigation.

### Conclusion

The **AI Path Planner** successfully integrates both **classical search algorithms** and **reinforcement learning** to demonstrate how different AI paradigms approach problem solving.

While  $A$  (*Manhattan*)<sup>\*</sup> remains the most optimal in static environments, the **Goal-Conditioned Q-Learning agent** showcases adaptive, efficient learning suited for dynamic or partially known worlds.

### Future Work:

- Incorporate **dynamic obstacle detection and real-time re-planning**.
- Extend the environment to **weighted terrains** (variable movement costs).
- Develop **reinforcement learning with continuous state spaces** and **policy gradient methods** for smoother navigation.
- Deploy a **hybrid planner** that switches between A\* and RL based on uncertainty or environment changes.

